ForecastingPath · Summary Report
2026-05-16T21:50:17.483206+00:00 · git 0b5f3dea · 26 resolved events

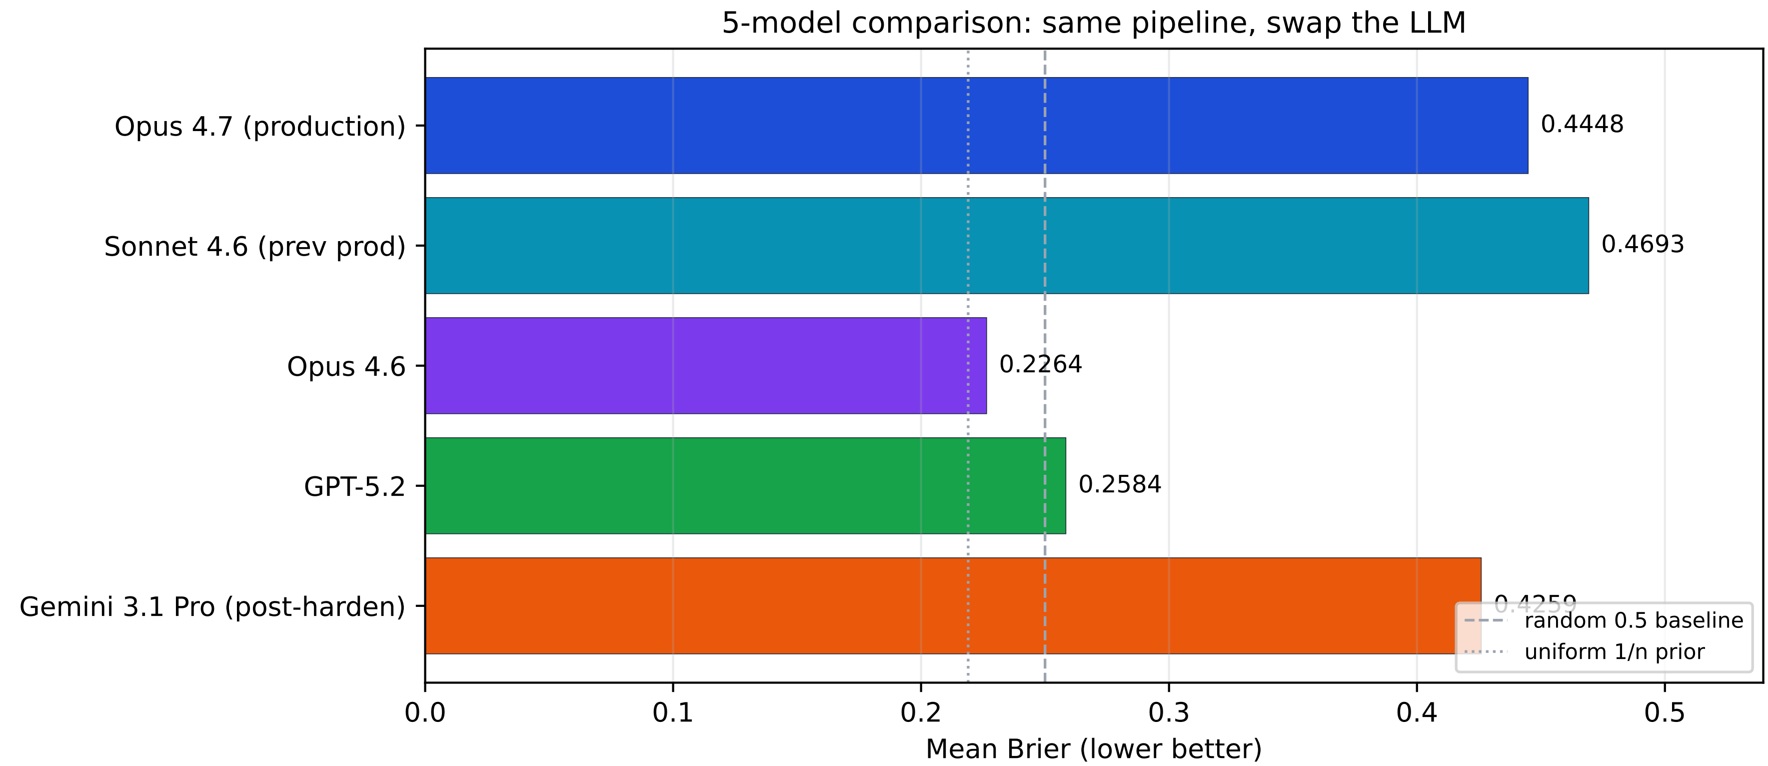

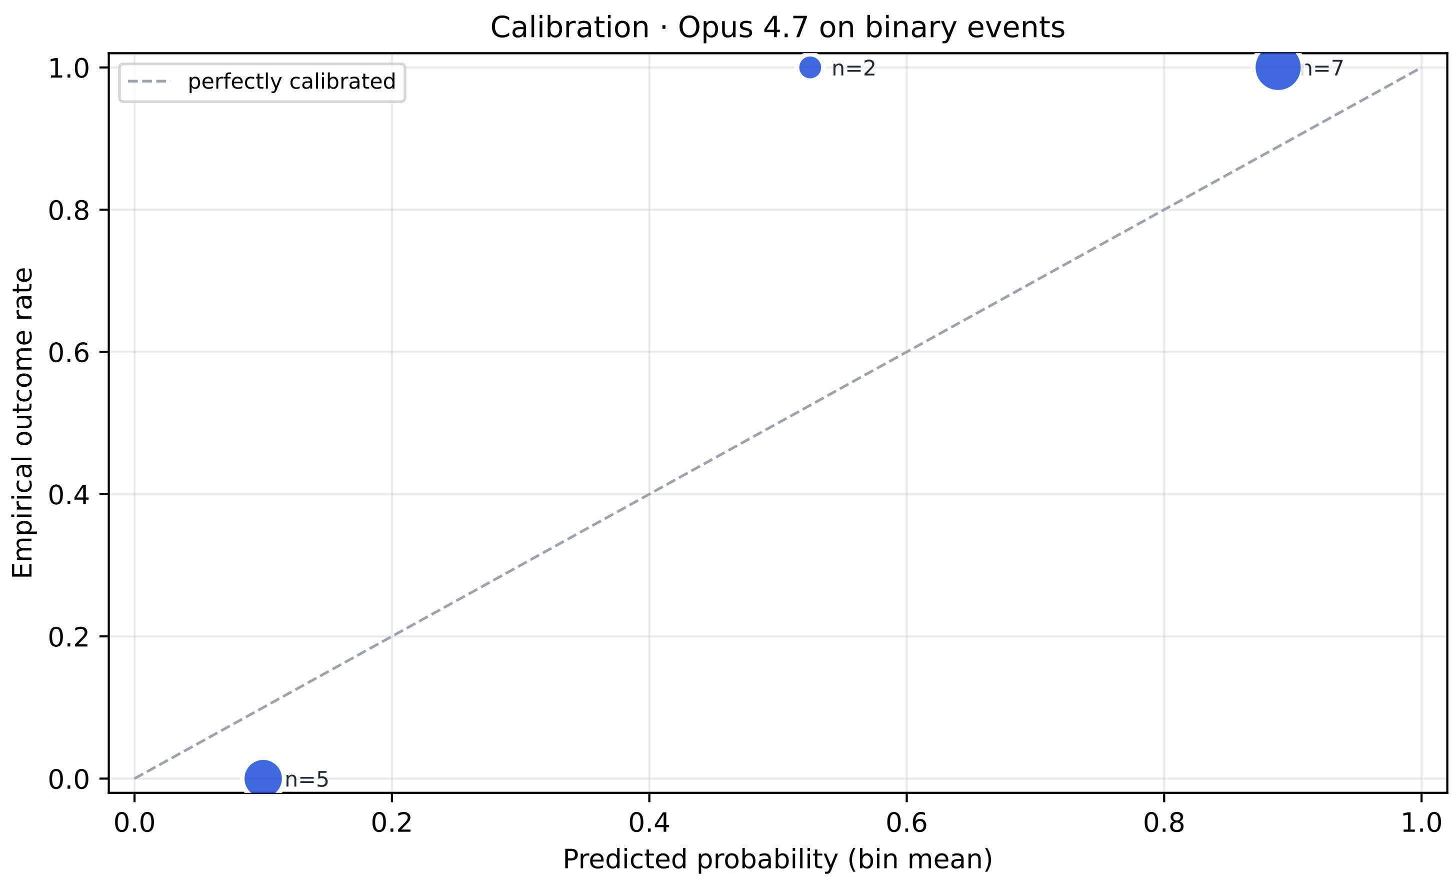

Per-event Brier + Findings

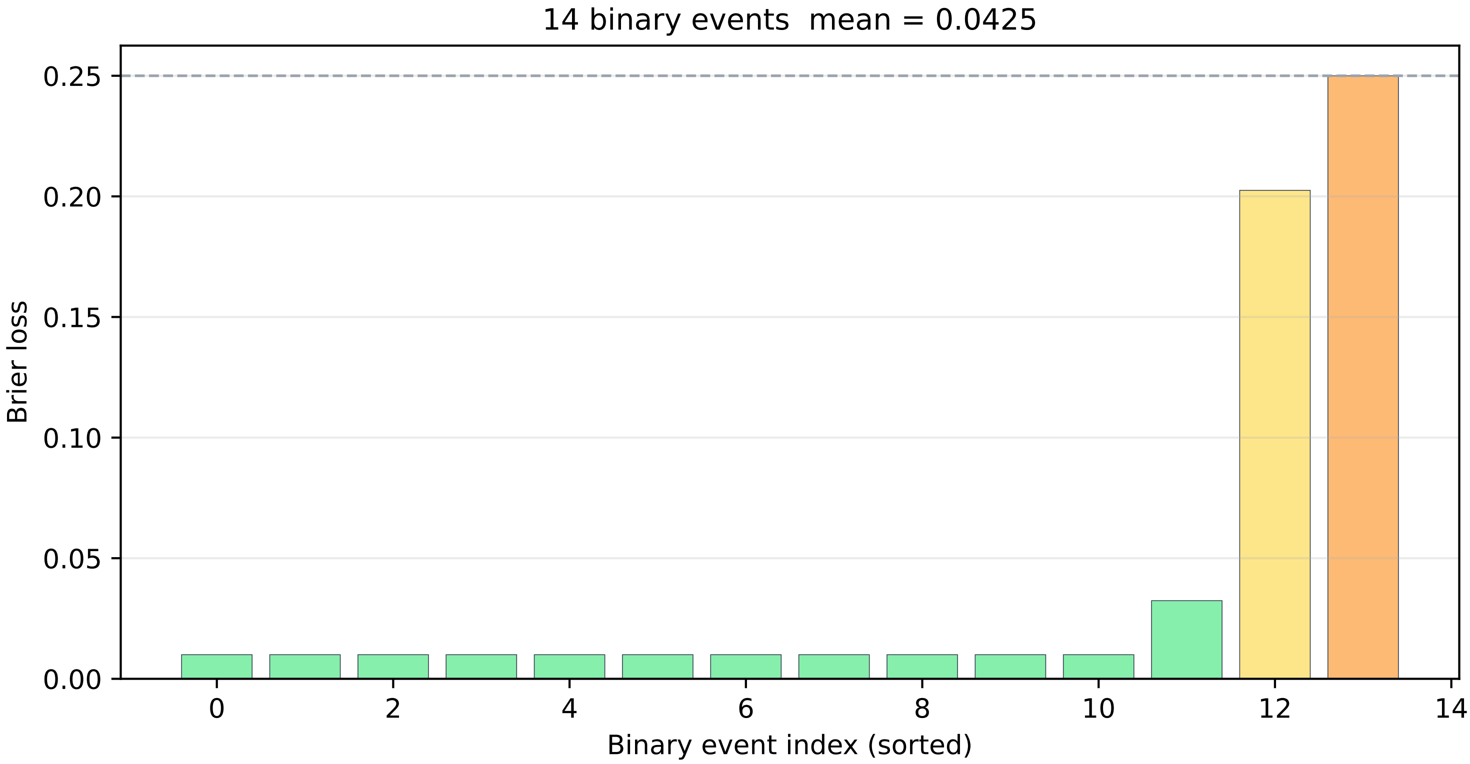

Findings worth keeping

1. The Opus 4.7 win is JSON schema compliance, not raw reasoning.
   Three of four alternative LLMs had Brier ≥ 0.22, dominated by
   catastrophic multi-outcome JSON failures.

2. Silent production bug accounted for most of the Phase 2 win.
   Old longshot floor max(0.05, 0.5/n) returned 0.25 for binary,
   silently clamping every binary prediction into [0.25, 0.75].
   New formula caps at the Kalshi-paper 0.10 threshold.
   ~6× per-event Brier improvement on binary longshots.

3. Leaderboard #1 ≠ best in your pipeline.
   Gemini 3.1 Pro Preview is public PA leaderboard's #1. In our
   pipeline it placed last with multi-outcome Brier 0.81.

4. Defensive engineering caught real bugs.
   Parser hardening (5 stages), fuzzy outcome matching, outcomes-
   missing safety net. Verify gate now loud after silently
   swallowing pytest failures for an entire session.

Honest caveats

• 26 events is a small sample, skewed toward binary tennis matches.
• Live Prophet Arena performance may differ by category mix.
• Multi-vendor evidence is the most generalizable finding here.