



5-model comparison: same pipeline, swap the LLM




Calibration · Opus 4.7 on binary events



Per-event Brier + Findings

14 binary events mean = 0.0425



Findings worth keeping

1. The Opus 4.7 win is JSON schema compliance, not raw reasoning. Three of four alternative LLMs had Brier ≥ 0.22 , dominated by catastrophic multi-outcome JSON failures.
2. Silent production bug accounted for most of the Phase 2 win. Old longshot floor $\max(0.05, 0.5/n)$ returned 0.25 for binary, silently clamping every binary prediction into $[0.25, 0.75]$. New formula caps at the Kalshi-paper 0.10 threshold. ~6x per-event Brier improvement on binary longshots.
3. Leaderboard #1 $\neq$ best in your pipeline. Gemini 3.1 Pro Preview is public PA leaderboard's #1. In our pipeline it placed last with multi-outcome Brier 0.81.
4. Defensive engineering caught real bugs. Parser hardening (5 stages), fuzzy outcome matching, outcomes-missing safety net. Verify gate now loud after silently swallowing pytest failures for an entire session.

Honest caveats

- 26 events is a small sample, skewed toward binary tennis matches.
- Live Prophet Arena performance may differ by category mix.
- Multi-vendor evidence is the most generalizable finding here.